

ARM Load/Store Instructions

- The ARM is a Load/Store Architecture:
 - Only load and store instructions can access memory
 - Does not support memory to memory data processing operations.
 - Must move data values into registers before using them.

ARM Load/Store Instructions

- ARM has three sets of instructions which interact with main memory. These are:
 - Single register data transfer (LDR/STR)
 - Block data transfer (LDM/STM)
 - Single Data Swap (SWP)

ARM Single Register Load/Store Instructions

- The basic load and store instructions are:

LDR	STR	Word
-----	-----	------

LDRB	STRB	Byte
------	------	------

LDRH	STRH	Halfword
------	------	----------

LDRSB		Signed byte load
-------	--	------------------

LDRSH		Signed halfword load
-------	--	----------------------

ARM Single Register Load/Store Instructions

- Memory system must support all access sizes
- Syntax:
 - LDR{<cond>}{<size>} Rd, <address>
 - STR{<cond>}{<size>} Rd, <address>

e.g.

LDR R0, [R1]

STR R0, [R1]

LDREQB R0, [R1]

Data Transfer: Memory to Register (load)

- To transfer a word of data, we need to specify two things:
 - Register: r0 - r15
 - Memory address: more difficult
 - Think of memory as a single one-dimensional array, so we can address it simply by supplying a pointer to a memory address.
 - There are times when we will want to offset from this pointer.

ARM Addressing Modes

There are basically two types of addressing modes available in ARM

- Pre-indexed addressing: the address generated is used immediately
- Post-indexed addressing: the address generated later replaces the base register

ARM Addressing Modes

[Rn]


Register

Address accessed is value found in Rn.

Example:

`ldr r0, [r1]`

@ $r0 \leftarrow *r1$




ARM Addressing Modes (Pre-Indexing)

[Rn, #±imm] *Immediate offset*

Address accessed is *imm* more/less than the address found in Rn. Rn does not change.

Example:

`ldr r2, [r1, #12]` @ $r2 \leftarrow *(r1 + 12)$



ARM Addressing Modes (Pre-Indexing)

$[Rn, \pm Rm]$

Register offset

Address accessed is the value in $Rn \pm$ the value in Rm . Rn and Rm do not change values.

Example:

`ldr r2, [r0, r1]` @ $r2 \leftarrow *(r0 + r1)$

ARM Addressing Modes (Pre-Indexing)

[Rn, \pm Rm, shift] *Scaled register offset*

Address accessed is the value in Rn \pm the value in Rm shifted as specified. Rn and Rm do not change values.

Example:

`ldr r0, [r1, r2, lsl #2]` @ $r0 \leftarrow *(r1 + r2 * 4)$

ARM Addressing Modes (Pre-Indexing w\ update)

$[Rn, \# \pm imm]!$

Immediate pre-indexed w\update

Address accessed is as with *immediate offset* mode, but Rn 's value updates to become the address accessed.

Example:

`ldr r2, [r1, #12]!` @ $r1 \leftarrow r1 + 12$ then $r2 \leftarrow *r1$

ARM Addressing Modes (Pre-Indexing w\ update)

$[Rn, \pm Rm]!$

Register pre-indexed w\update

Address accessed is as with *register offset* mode, but Rn 's value updates to become the address accessed.

Example:

`ldr r2, [r0, r1]!` @ $r0 \leftarrow r0 + r1$ then $r2 \leftarrow *r0$

ARM Addressing Modes (Pre-Indexing w\ update

$[Rn, \pm Rm, shift]!$ *Scaled register pre-indexed w\update*
Address accessed is as with *scaled register offset* mode, but Rn 's value updates to become the address accessed.

Example:

`ldr r2, [r0, r1, lsl #2]!` @ $r0 \leftarrow r0 + r1 * 4$ then $r2 \leftarrow *r0$

ARM Addressing Modes (Post-Indexing)

$[Rn], \# \pm imm$

Immediate post-indexed

Address accessed is value found in Rn , and then Rn 's value is increased/decreased by imm .

Example:

`str r2, [r1], +4`

@ $*r1 \leftarrow r2$ then $r1 \leftarrow r1 + 4$

ARM Addressing Modes (Post-Indexing)

$[Rn], \pm Rm$

Register post-indexed

Address accessed is value found in Rn , and then Rn 's value is increased/decreased by Rm .

Example:

`str r0, [r1], r2` @ $*r1 \leftarrow r0$ then $r1 \leftarrow r1 + r2$

ARM Addressing Modes (Post-Indexing)

$[Rn], \pm Rm, shift$ *Scaled register post-indexed*

Address accessed is value found in Rn , and then Rn 's value is increased/decreased by Rm shifted according to *shift*.

Example:

`ldr r0, [r1], r2, lsl #3` @ $r0 \leftarrow *r1$ then $r1 \leftarrow r1 + r2 * 8$

Examples of pre- and post-indexed addressing

str r3, [r0, r4, lsl #3]

@ *pre-indexed*

ldr r5, [r0, r1, lsl #3]!

@ pre indexed with writeback

ldr r0, [r1, #-8]

@ pre-indexed with negative offset

ldr r0, [r1, -r2, lsl #2]

@ negative offset shifted

ldrb r5, [r1]

@ load byte from ea <r1>

ldrsh r5, [r3]

@ load signed halfword from ea <r3>

ldrsb r5, [r3, #0xc1]

@ load signed byte from ea <r3+193>

str r7, [r0], #4

@ store r7 to ea<r0>, then add #24 to r0

ldr r2, [r0], r4, lsl #2

@ load r2 from ea<r0>, then add r4*4 to r0

ldrh r3, [r5], #2

@ load halfword to r3 from ea<r5>, then

@ add #2 to r5

strh r2, [r5], #8

@ store halfword from r2 to ea<r5>, then

@ add 8 to r5

Block Data (Multiple data) transfer instructions

- ARM also supports multiple loads and stores:
- When the data to be copied to the stack is known to be a multiple of 4 bytes, copying is faster using load multiple and store multiple instructions.
- Each instruction can transfer 1 or more 32-bit words between registers and memory
- Each instruction updates another register that holds the memory address

Multiple data transfer instructions

General syntax:

`op<address-mode>{cond} <rn>{!}, <register-list>{^}`

- `op` : `ldm`, `stm`
- `address-mode`:

`ia` – **Increment** address **after** each transfer

`ib` – **Increment** address **before** each transfer.

`da` – **Decrement** address **after** each transfer

`db` – **Decrement** address **before** each transfer

`fd` – **full descending** stack

`ed` – **empty** descending stack

`fa` – **full ascending** stack

`ea` – **empty ascending** stack.

Multiple data transfer instructions

- *cond* is an optional condition code
- *rn* is the *base register* containing the **initial memory address** for the transfer.
- *!* is an optional suffix.
 - If *!* is present, the final address is written back into *rn*.
 - If the base register is in the register-list, then you must not use the writeback option.

Multiple data transfer instructions

reg-list

- a list of registers to be loaded or stored.
- can be a comma-separated list or an rx-ry range.
- may contain any or all of r0 - r15
- the registers are **always** loaded in order, regardless to how the registers are ordered in the list.
- for both the ldm and stm instructions, reg-list must not contain the sp
- for ldm, reg-list must not contain the PC if it contains the lr
- for stm, reg-list must not contain the lr if it contains the pc

Multiple data transfer instructions

- \wedge is an optional suffix. Do NOT use it in User mode or System mode.
 - forces processor to transfer the saved program status register (SPSR) into the current program status register (CPSR) at the same time, saving us an instruction
 - if *op* is LDM and *register-list* contains the pc, the CPSR is restored from the SPSR
 - otherwise, data is transferred into or out of the User mode registers instead of the current mode registers.

Multiple data transfer instructions

Instruction	Operation	[s]	Notes
ldmia $r_n[!], <reglist>$	$r_{first} \leftarrow \text{mem}[r_n]$ $r_{first+1} \leftarrow \text{mem}[r_n+4]$ $r_{first+2} \leftarrow \text{mem}[r_n+8]$ etc.		$r_{first}, r_{last} \dots$ refer to the order of the registers by register number, not by the order they appear in the list '!' is an optional write-back suffix. If present, the final address that is written to or stored from is written back into r_n
ldmdb $r_{last}[!], <reglist>$	$r_{last} \leftarrow \text{mem}[r_n]$ $r_{last-1} \leftarrow \text{mem}[r_n-4]$ $r_{last-2} \leftarrow \text{mem}[r_n-8]$ etc.		
stmia $r_n[!], <reglist>$	$\text{mem}[r_n] \leftarrow r_{first}$ $\text{mem}[r_n+4] \leftarrow r_{first+1}$ $\text{mem}[r_n+8] \leftarrow r_{first+2}$ etc.		
stmdb $r_n[!], <reglist>$	$\text{mem}[r_n] \leftarrow r_{last}$ $\text{mem}[r_n-4] \leftarrow r_{last-1}$ $\text{mem}[r_n-8] \leftarrow r_{last-2}$ etc.		

Multiple data transfer instructions

Example of ldmia – load, increment after

```
ldmia    r9, {r0-r3}    // register 9 holds the  
                        // base address. "ia" says  
                        // increment the base addr  
                        // after each value has  
                        // been loaded from memory
```


Multiple data transfer instructions

Example of `ldmia` – load, increment after

```
ldmia  r9, {r0-r3}    // register 9 holds the
                        // base address
```

This has the same effect as four separate `ldr` instructions, or

```
ldr    r0, [r9]
ldr    r1, [r9, #4]
ldr    r2, [r9, #8]
ldr    r3, [r9, #12]
```

Note: at the end of the `ldmia` instruction, register `r9` has not been changed. If you wanted to change `r9`, you could simply use

```
ldmia  r9!, {r0-r3, r12}
```

Multiple register data transfer instructions

ldmia – Example 2

```
ldmia r9, {r0-r3, r12}
```

- Load words addressed by r9 into r0, r1, r2, r3, and r12
- Increment r9 after each load.

Example 3

```
ldmia r9, {r5, r3, r0-r2, r14}
```

- load words addressed by r9 into registers r0, r1, r2, r3, r5, and r14.
- Increment r9 after each load.
- ldmib, ldmda, ldmdb work similar to ldmia
- Stores work in an analogous manner to load instructions

PUSH and POP

Note:

push is a synonym for stmdb sp!, reg-list

pop is a synonym for ldmia sp!, reg-list

Note:

ldmfd is a synonym for ldmia

stmfd is a synonym for stmdb

Multiple register data transfer instructions

Common usage of multiple data transfer instructions

- **Stack**
 - Function calls
 - Context switches
 - Exception handlers

Multiple register data transfer instructions

Stack

- When making nested subroutine calls, we need to store the current state of the processor.
- The multiple data transfer instructions provide a mechanism for storing state on the *runtime stack* (pointed to by the stack pointer, r13 or sp)

stack addressing:

- stacks can *ascend* or *descend* memory
- stacks can be *full* or *empty*
- ARM multiple register transfers support all forms of the stack

Multiple register data transfer instructions

Stack

- Ascending stack: grows **up**
- Descending stack: grows **down**

A **stack pointer (sp)** holds the address of the current top of the stack

Full stack: sp is pointing to the last valid data item pushed onto the stack

Empty stack: sp is pointing to the vacant slot where the next data item will be placed

Multiple register data transfer instructions

Stack Processing

ARM support for all four forms of stacks

- *Full ascending (FA)*: grows up; stack pointer points to the highest address containing a valid data item
- *Empty ascending (EA)*: grows up; stack pointer points to the first empty location
- *Full descending (FD)*: grows down; stack pointer points to the lowest address containing a valid data item
- *Empty descending (ED)*: grows down; stack pointer points to the first empty location below the stack

Load and Store Multiples

LDMxx r10, {r0,r1,r4}

STMxx r10, {r0,r1,r4}

Base Register (Rb)

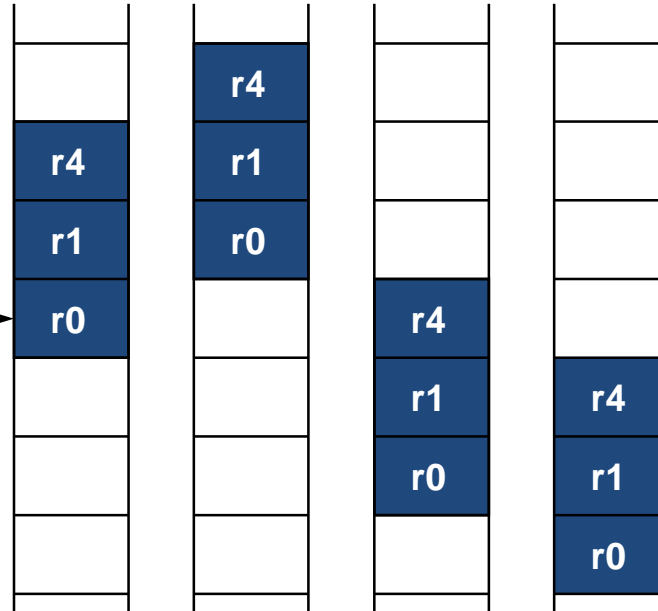
r10

IA

IB

DA

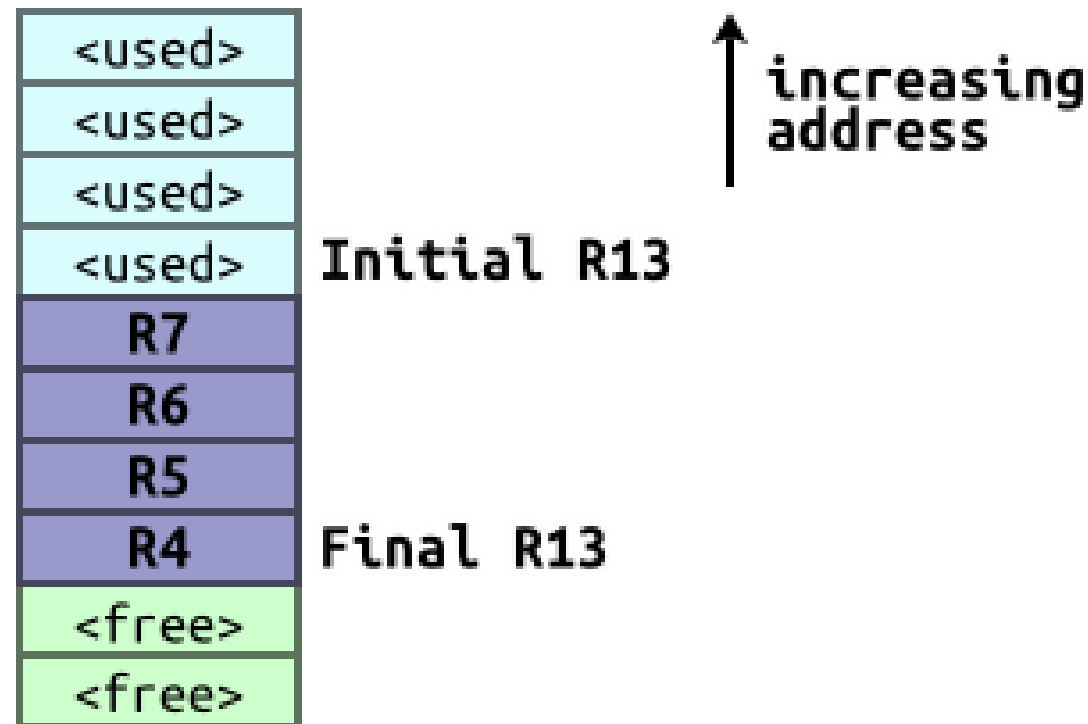
DB



Increasing
Address

Stack push operation: stmfd

STMFD r13!, {r4-r7} – Push R4,R5,R6 and R7 onto the stack.



Stack pop operation: ldmed

LDMFD r13!, {r4-r7} – Pop R4,R5,R6 and R7 from the stack.

